


# Postgres

## The Best Tool You're Already Using

- [Adam Sanderson](#)
- [LiquidPlanner](#)

≈ **LiquidPlanner**



# Adam Sanderson

I have been a full stack engineer at LiquidPlanner for 5 years.

- I got off in Kansas\*, and that's ok!
- Github: [adamsanderson](#)
- Twitter: [adamsanderson](#)
- Blog: <http://monkeyandcrow.com>

\* Seattle

# ≈ LiquidPlanner®

Online project management with probabilistic scheduling.

- Started in 2007 with Rails 1.x
- Used Postgres from the beginning
- We have learned some great techniques along the way

# Topics

- Tagging
- Hierarchy
- Custom Data
- Full Text Search

# Method

For each topic, we'll cover the SQL before we cover its use in ActiveRecord.

We will use Postgres 9.x, Ruby 1.9 syntax, and ActiveRecord 4.0.

If you understand the SQL you can use it in any version of ActiveRecord, 4.0 just makes it easier.

# Backstory

You just built a great new social network for hedgehog lovers around the world,  
[HedgeWith.me](http://HedgeWith.me).

Everything is going well. You have a few users, but now they want more.



*My hedgehog is afraid of grumpy hedgehogs, but likes cute ones how can I find him friends?*

***hedgehogs4life***

## Tagging

People want to be able to tag their hedgehogs, and then find other hedgehogs with certain tags.



# Defining Arrays in SQL

```
CREATE TABLE hedgehogs (  
    id      integer primary key,  
    name    text,  
    age     integer,  
    tags    text[ ]  
);
```



# Defining Arrays in ActiveRecord

```
create_table :hedgehogs do |t|  
  t.string :name  
  t.integer :age  
  t.text :tags, array: true  
end
```

ActiveRecord 4.x introduced arrays for Postgres, use `array: true`

# Heads Up

Define array columns as `t.text` instead of `t.string` to avoid casting.

Postgres assumes that `ARRAY['cute', 'cuddly']` is of type `text[]` and will require you to cast, otherwise you will see errors like this:

```
ERROR: operator does not exist: character varying[] && text[]
```

# Boolean Set Operators

You can use the set operators to query arrays.

- `A @> B` A contains all of B
- `A && B` A overlaps any of B

# Querying Tags in SQL

Find all the hedgehogs that are **spiny** or **prickly**:

```
SELECT name, tags FROM hedgehogs  
WHERE tags && ARRAY['spiny', 'prickly'];
```

A && B    A overlaps any of B

# Querying Tags in SQL

<b>name</b>	<b>tags</b>
Marty	<code>spiny</code> , <code>prickly</code> , cute
Quilby	cuddly, <code>prickly</code> , hungry
Thomas	grumpy, <code>prickly</code> , sleepy, <code>spiny</code>
Franklin	<code>spiny</code> , round, tiny

# Querying Tags in SQL

Find all the hedgehogs that are **spiny** and **prickly**:

```
SELECT name, tags FROM hedgehogs  
WHERE tags @> ARRAY['spiny', 'prickly'];
```

A @> B    A contains all the B

# Querying Tags in SQL

name	tags
Marty	<code>spiny</code> , <code>prickly</code> , cute
Thomas	grumpy, <code>prickly</code> , sleepy, <code>spiny</code>

# Querying Tags in ActiveRecord

Find all the hedgehogs that are spiny and prickly

```
Hedgehog.where "tags @> ARRAY[?]", ['spiny', 'prickly']
```



# Querying Tags in ActiveRecord

Create scopes to encapsulate set operations:

```
class Hedgehog < ActiveRecord::Base  
  scope :any_tags, -> (* tags){where('tags && ARRAY[?]', tags)}  
  scope :all_tags, -> (* tags){where('tags @> ARRAY[?]', tags)}  
end
```

# Querying Tags in ActiveRecord

Find all the hedgehogs that are spiny or large, and older than 4:


```
Hedgehog.any_tags('spiny', 'large').where('age > ?', 4)
```

*Hi, I run an influential hedgehog club. Our members would all use [HedgeWith.me](https://hedgewith.me), if they could show which hogs are members of our selective society.*

***Boston Spine Fancy President***

## Hierarchy

Apparently there are thousands of hedgehog leagues, divisions, societies, clubs, and so forth.



# Hierarchy

We need to efficiently model a club hierarchy like this:

- North American League
  - Western Division
    - Cascadia Hog Friends
    - Californian Hedge Society

How can we support operations like finding a club's depth, children, or parents?

# Materialized Path in SQL

Encode the parent ids of each record in its `path`.

```
CREATE TABLE clubs (  
    id            integer primary key,  
    name          text,  
    path          integer[ ]  
);
```

# Querying a Materialized Path

<b>id</b>	<b>name</b>	<b>path</b>
1	North American League	[1]
2	Eastern Division	[1,2]
4	New York Quillers	[1,2,4]
5	Boston Spine Fancy	[1,2,5]
3	Western Division	[1,3]
6	Cascadia Hog Friends	[1,3,6]
7	California Hedge Society	[1,3,7]
...		

# Materialized Path: Depth

The depth of each club is simply the length of its path.

- `array_length(array, dim)` returns the length of the array

`dim` will always be 1 unless you are using multidimensional arrays.

# Materialized Path: Depth

Display the top two tiers of hedgehog clubs:

```
SELECT name, path, array_length(path, 1) AS depth
FROM clubs
WHERE array_length(path, 1) <= 2
ORDER BY path;
```

`array_length(path, 1)` is the depth of record



# Materialized Path: Depth

<b>name</b>	<b>path</b>	<b>depth</b>
North American League	[1]	1
Eastern Division	[1,2]	2
Western Division	[1,3]	2
South American League	[9]	1

# Materialized Path: Children

Find all the clubs that are children of the California Hedge Society, ID: 7.

```
SELECT id, name, path FROM clubs
WHERE path && ARRAY[7]
ORDER BY path
```

A && B    A overlaps any of B

# Materialized Path: Children

id	name	path
7	Californian Hedge Society	[1,3,7]
8	Real Hogs of the OC	[1,3,7,8]
12	Hipster Hogs	[1,3,7,12]

Apparently it is illegal to own hedgehogs in California

# Materialized Path: Parents

Find the parents of the California Hedge Society, Path: `ARRAY[1,3,7]` .

```
SELECT name, path FROM clubs
WHERE ARRAY[id] && ARRAY[1,3,7]
ORDER BY path;
```

`A && B` A overlaps any of B

# Materialized Path: Parents

<b>id</b>	<b>name</b>	<b>path</b>
1	North American League	[1]
3	Western Division	[1,3]
7	Californian Hedge Society	[1,3,7]

# ActiveRecord: Arrays & Depth

With ActiveRecord 4.x, `path` is just ruby array.

```
class Club < ActiveRecord::Base
  def depth
    self.path.length
  end
  ...
end
```

# Querying in ActiveRecord

Encapsulate these conditions as instance methods:

```
class Club < ActiveRecord::Base

  def children

    Club.where('path && ARRAY[?]', self.id)

  end

  def parents

    Club.where('ARRAY[id] && ARRAY[?]', self.path)

  end

end
```

# Querying in ActiveRecord

Now we have an easy way to query the hierarchy.

```
@club.parents.limit(5)
```

```
@club.children.joins(:hedgehogs).merge(Hedgehog.any_tags('silly'))
```

These features can all work together.

Mind blown?



*I need to keep track of my hedgehogs' favorite foods, colors, weight, eye color, and shoe sizes!*

### ***the Quantified Hedgehog Owner***

*If I am forced to enter my hedgehog's shoe size, I will quit immediately!*

### ***the Unquantified Hedgehog Owner***

## Custom Data



Your users want to record arbitrary data about their hedgehogs.

# Hstore

Hstore provides a hash column type. It is a useful alternative to ActiveRecord's `serialize` where the keys and values can be queried in Postgres.

# Hstore

Hstore needs to be installed manually. Your migration will look like this:

```
class InstallHstore < ActiveRecord::Migration
  def up
    execute 'CREATE EXTENSION hstore'
  end
  ...
end
```

# Heads Up

Although hstore is supported by ActiveRecord 4.x, the default schema format does not support extensions.

Update `config/application.rb` to use the SQL schema format, otherwise your tests will fail.

```
class Application < Rails::Application
  config.active_record.schema_format = :sql
end
```

# Defining an Hstore in SQL

```
CREATE TABLE hedgehogs (  
    id          integer primary key,  
    name       text,  
    age        integer,  
    tags       text[],  
    custom     hstore DEFAULT '' NOT NULL  
);
```

# Defining an Hstore in ActiveRecord

`hstore` is supported in ActiveRecord 4.x as a normal column type:

```
create_table :hedgehogs do |t|  
  t.string :name  
  t.integer :age  
  t.text :tags, array: true  
  t.hstore :custom, :default => '', :null => false  
end
```

# Heads Up

Save yourself some hassle, and specify an empty hstore by default:

```
t.hstore :custom, :default => '', :null => false
```

Otherwise new records will have null hstores.

# Hstore Format

Hstore uses a text format, it looks a lot like a ruby 1.8 hash:

```
UPDATE hedgehogs SET
```

```
custom = '"favorite_food" => "lemons", "weight" => "2lbs"'
```

```
WHERE id = 1;
```

Be careful of quoting.



# Hstore Operators

Common functions and operators:

- `defined(A, B)` Does A have B?
- `A -> B` Get B from A. In ruby this would be `A[B]`

# Query Hstore in SQL

Find all the favorite foods of the hedgehogs:

```
SELECT name, custom -> 'favorite_food' AS food  
FROM hedgehogs WHERE defined(custom, 'favorite_food');
```

`defined(A, B)` Does A have B?

`A -> B` Get B from A. In ruby this would be `A[B]`

# Query Hstore in SQL

**name**

Horrace

Quilby

Thomas

**food**

lemons

pasta

grubs

# Query Hstore in ActiveRecord

Create scopes to make querying easier:

```
class Hedgehog < ActiveRecord::Base

  scope :has_key,    -> (key){ where('defined(custom, ?)', key) }
  scope :has_value, -> (key, value){ where('custom -> ? = ?', key, value) }
  ...
end
```

# Query Hstore in ActiveRecord

Find hedgehogs with a custom `color`:

```
Hedgehog.has_key( 'color' )
```

# Query Hstore in ActiveRecord

Find hedgehogs that are brown:

```
Hedgehog.has_value('color', 'brown')
```

# Query Hstore in ActiveRecord

Find all the silly, brown, hedgehogs:

```
Hedgehog.any_tags('silly').has_value('color', 'brown')
```

# Updating an Hstore with ActiveRecord

With ActiveRecord 4.x, hstore columns are just hashes:

```
hedgehog.custom["favorite_color"] = "ochre"
```

```
hedgehog.custom = {favorite_food: "Peanuts", shoe_size: 3}
```



# Heads Up

Hstore columns are always stored as strings:

```
hedgehog.custom["weight"] = 3
```

```
hedgehog.save!
```

```
hedgehog.reload
```

```
hedgehog.custom['weight'].class #=> String
```

*Someone commented on my hedgehog. They said they enjoy his beady little eyes, but I can't find it.*

***hogmama73***

## Full Text Search

Your users want to be able to search within their comments.



# Full Text Search in SQL

```
CREATE TABLE comments (  
    id                integer primary key,  
    hedgehog_id       integer,  
    body              text  
);
```

# Full Text Search Data Types

There are two important data types:

- `tsvector` represents the text to be searched
- `tsquery` represents the search query

# Full Text Search Functions

There are two main functions that convert strings into these types:

- `to_tsvector(configuration, text)` creates a normalized `tsvector`
- `to_tsquery(configuration, text)` creates a normalized `tsquery`

# Full Text Search Normalization

Postgres removes common stop words:

```
select to_tsvector('A boy and his hedgehog went to Portland');
```

```
-- boy, hedgehog, portland, went
```

```
select to_tsvector('I need a second line to fill space here.');
```

```
-- fill, line, need, second, space
```

# Full Text Search Normalization

Stemming removes common endings from words:

## **term**

hedgehogs

enjoying

piping

## **stemmed**

hedgehog

enjoy

pipe

# Full Text Search Operators

## Vectors:

- `v @@ Q` Searches V for Q

## Queries:

- `v @@ (A && B)` Searches V for A and B
- `v @@ (A || B)` Searches V for A or B



# Full Text Search Querying

Find comments about "enjoying" something:

```
SELECT body
FROM comments
WHERE to_tsvector('english', body)
    @@ to_tsquery('english', 'enjoying');
```

v @@ Q Searches V for Q

# Full Text Search Querying

- Does he **enjoy** beets? Mine loves them
- I really **enjoy** oranges
- I am **enjoying** these photos of your hedgehog's beady little eyes
- Can I feed him grapes? I think he **enjoys** them.

Notice how "enjoying" also matched "enjoy" and "enjoys" due to stemming.

# Full Text Search Wildcards

- `to_tsquery('english','cat:*')` Searches for anything starting with cat

Such as: **cat**, **cat**apult, **cat**aclysmic.

But not: octo**cat**, s**cat**ter, prognosti**cate**

# Full Text Search Wild Cards

Find comments containing the term "oil", and a word starting with "quil" :

```
SELECT body
FROM comments
WHERE to_tsvector('english', body)
    @@ ( to_tsquery('english','oil')
        && to_tsquery('english','quil:*')
    );
```

`V @@ (A && B)` Searches V for A and B

# Full Text Search Querying

- What brand of oil do you use? Have you tried QuillSwill?

# Heads Up

`tsquery` only supports wildcards at the end of a term.

While `quill:*` will match "QuillSwill", but `*:swill` will not.

In fact, `*:swill` will throw an error.

# Even More Heads Up!

Never pass user input directly to `to_tsquery`, it has a strict mini search syntax. The following all fail:

- `http://localhost` `:` has a special meaning
- `O'Reilly's Books` Paired quotes cannot be in the middle
- `A && B` `&` and `|` are used for combining terms

You need to sanitize queries, or use a gem that does this for you.

# Full Text Search With ActiveRecord

We can wrap this up in a scope.

```
class Comment < ActiveRecord::Base  
  scope :search_all, -> (query){  
    where("to_tsvector('english', body) @@ #{sanitize_query(query)}")  
  }  
end
```

You need to write `sanitize_query`, or use a gem that does this for you.



# Full Text Search With ActiveRecord

Find the comments about quill oil again, and limit it to 5 results:

```
Comment.search_all("quill* oil").limit(5)
```

Since `search_all` is a scope, we chain it like all the other examples.

# Full Text Search Indexing

Create an index on the function call `to_tsvector('english', body)`:

```
CREATE INDEX comments_gin_index
ON comments
USING gin(to_tsvector('english', body));
```

The `gin` index is a special index for multivalued columns like a `text[]` or a `tsvector`

# Heads Up

Since we are indexing a function call, `to_tsvector('english', body)`, we must call it the same way every time.

You don't have to use `english`, but you do need to be consistent.

# In Summary

- Arrays can model tagging and hierarchies
- Hstore can be used to model custom data
- Postgres supports full text search

You can now enjoy the happy hour!

```
SELECT * FROM beers WHERE  
traits @> ARRAY[ 'hoppy' , 'floral' ]
```

# Any Questions?

Possible suggestions:

- Why not normalize your database instead of using arrays?
- Can I see how you implemented `sanitize_query`?
- What is a good gem for full text search?
- What about ActiveRecord 2 and 3?
- Why hstore instead of JSON?
- Can I buy you coffee?

# Extra Resources

- ActiveRecord [Queries & Scopes](#)
- Postgres [Array Operators](#)
- Postgres [Hstore Documentation](#)
- Postgres [Full Text Search](#)
- Ruby Gems for Full Text Search
  - [Textacular](#) Supports Active Record 2.x and 3.x
  - [pg\\_search](#) Supports Active Record 3.x, but has more features
- My [Blog](#), [Github](#), and [favorite social network](#)
- [How to draw a hedgehog.](#)

# Bonus

Here's `sanitize_query`:

```
def self.sanitize_query(query, conjunction=' && ')  
  "(" + tokenize_query(query).map{|t| term(t)}.join(conjunction) + ")"  
end
```

It breaks up the user's request into terms, and then joins them together.

# Bonus

We tokenize by splitting on white space, `&`, `|`, and `:`.

```
def self.tokenize_query(query)
  query.split(/(\s|[\&|:])+/)
end
```



# Bonus

Each of those tokens gets rewritten:

```
def self.term(t)
  # Strip leading apostrophes, they are never legal, "'ok" becomes "ok"
  t = t.gsub(/^'+/, '')

  # Strip any *s that are not at the end of the term
  t = t.gsub(/\*[^\$]/, '')

  # Rewrite "sear*" as "sear:*" to support wildcard matching on terms
  t = t.gsub(/\*$/, ':*')

  ...
```

...

*# If the only remaining text is a wildcard, return an empty string*

```
t = "" if t.match(/^[:* ]+$/)
```

```
"to_tsquery('english', #{quote_value t})"
```

```
end
```